

# Part 1: Data Acquisition & Preprocessing

**Objective:** Acquire and preprocess S&P 500 stock price data for portfolio analysis.

**Key Tasks:** 1. Scrape top 100 S&P 500 tickers by market capitalization 2. Download 5 years of End-of-Day OHLCV data 3. Clean and align data to uniform trading calendar 4. Compute daily log returns

**Deliverables:** - **prices:** Adjusted close price matrix (DataFrame) - **log\_returns:** Daily log returns matrix (DataFrame) - Complete code workflow for data acquisition and preprocessing

```
# Import required libraries
import yfinance as yf
import pandas as pd
import numpy as np
import warnings
import matplotlib.pyplot as plt
warnings.filterwarnings('ignore')

print("Libraries imported successfully!")
```

Libraries imported successfully!

## Task 1: Scrape Top 100 S&P 500 Tickers by Market Cap

```
# Get S&P 500 tickers from Wikipedia and select top 100 by market cap
def get_sp500_tickers():
    """Scrape S&P 500 tickers from Wikipedia"""
    url = 'https://en.wikipedia.org/wiki/List_of_S%26P_500_companies'
    tables = pd.read_html(url)
    sp500_table = tables[0]
```

```

    tickers = sp500_table['Symbol'].str.replace('.', '-').tolist()
    return tickers

# Get all S&P 500 tickers
sp500_tickers = get_sp500_tickers()
print(f"Found {len(sp500_tickers)} S&P 500 tickers")

# Get market caps to select top 100
print("Getting market capitalizations...")
market_caps = {}

# Process efficiently - check first 150 tickers
for ticker in sp500_tickers[:150]:
    try:
        stock = yf.Ticker(ticker)
        info = stock.info
        market_cap = info.get('marketCap', 0)
        if market_cap > 0:
            market_caps[ticker] = market_cap
    except:
        continue

# Sort by market cap and get top 100
sorted_tickers = sorted(market_caps.items(), key=lambda x: x[1], reverse=True)
top_100_tickers = [ticker for ticker, _ in sorted_tickers[:100]]

print(f" Selected top 100 tickers by market cap")
print(f"Sample tickers: {top_100_tickers[:5]}..." +
      f"{top_100_tickers[-5:]}")
print(f"Total: {len(top_100_tickers)} tickers")

```

Found 503 S&P 500 tickers

Getting market capitalizations...

Selected top 100 tickers by market cap

Sample tickers: ['AAPL', 'GOOG', 'GOOGL', 'AMZN', 'AVGO']...['STZ', 'AWK', 'AEE', 'ADM', 'AVI

Total: 100 tickers

Selected top 100 tickers by market cap

Sample tickers: ['AAPL', 'GOOG', 'GOOGL', 'AMZN', 'AVGO']...['STZ', 'AWK', 'AEE', 'ADM', 'AVI

Total: 100 tickers

## Task 2: Download EOD OHLCV Data (5 Years)

```
# Download 5 years of EOD OHLCV data
from datetime import datetime, timedelta

# Define date range
end_date = datetime.now()
start_date = end_date - timedelta(days=5*365)

print(f"Downloading data from {start_date.strftime('%Y-%m-%d')} to " +
      f"{end_date.strftime('%Y-%m-%d')} for {len(top_100_tickers)} tickers...")
# Download data for all tickers
raw_data = yf.download(
    tickers=top_100_tickers,
    start=start_date.strftime('%Y-%m-%d'),
    end=end_date.strftime('%Y-%m-%d'),
    group_by='ticker',
    auto_adjust=True,
    prepost=False,
    threads=True
)

print(f" Download completed!")
print(f"Data shape: {raw_data.shape}")
print(f"Date range: {raw_data.index[0]} to {raw_data.index[-1]}")

# Display first 5 rows of raw data
print("Displaying first 5 rows of raw data:")
print(raw_data.head(5))
```

Downloading data from 2020-08-07 to 2025-08-06 for 100 tickers...

[\*\*\*\*\*100%\*\*\*\*\*] 100 of 100 completed

```
Download completed!
Data shape: (1254, 500)
Date range: 2020-08-07 00:00:00 to 2025-08-05 00:00:00
Displaying first 5 rows of raw data:
Ticker      DXCM
Price      Open      High      Low      Close      Volume \
```

Date					
2020-08-07	110.220001	114.057503	108.754997	110.175003	3715200
2020-08-10	109.827499	109.957497	101.552498	105.305000	4454800
2020-08-11	103.752502	104.802498	100.787498	102.787498	3273200
2020-08-12	103.050003	107.605003	103.050003	106.972504	2465200
2020-08-13	108.199997	109.997498	107.684998	108.820000	2421600

Ticker	BR					...	\
Price	Open	High	Low	Close	Volume	...	
Date						...	
2020-08-07	124.052006	125.266405	123.555215	124.429207	768200	...	
2020-08-10	124.879978	126.085168	124.475178	125.671173	721200	...	
2020-08-11	127.722747	132.939111	127.639951	130.740326	1131400	...	
2020-08-12	130.363149	130.979544	127.777968	127.989563	837400	...	
2020-08-13	127.290356	129.029144	127.078761	127.731949	655400	...	

Ticker	BA						\
Price	Open	High	Low	Close	Volume		
Date							
2020-08-07	171.500000	171.860001	168.699997	170.020004	19318000		
2020-08-10	171.360001	179.789993	171.330002	179.410004	35857700		
2020-08-11	184.509995	189.970001	179.529999	180.130005	61036600		
2020-08-12	184.009995	184.149994	173.190002	175.440002	40674400		
2020-08-13	173.619995	179.470001	172.429993	174.729996	22958800		

Ticker	AZO					
Price	Open	High	Low	Close	Volume	
Date						
2020-08-07	1178.369995	1184.949951	1170.219971	1182.219971	200700	
2020-08-10	1185.260010	1186.880005	1169.140015	1174.719971	128600	
2020-08-11	1179.199951	1182.869995	1160.780029	1171.560059	162700	
2020-08-12	1176.849976	1193.790039	1174.449951	1188.369995	167000	
2020-08-13	1183.359985	1190.920044	1177.310059	1187.140015	103900	

[5 rows x 500 columns]

### Task 3: Clean and Align Data to Uniform Calendar

```

# Extract and clean adjusted close prices
print("Extracting and cleaning adjusted close prices...")

# Extract close prices for all tickers
if len(top_100_tickers) == 1:
    # Single ticker case
    prices = raw_data[['Close']].copy()
    prices.columns = top_100_tickers
else:
    # Multiple tickers case
    prices = pd.DataFrame(index=raw_data.index)
    for ticker in top_100_tickers:
        if ticker in raw_data.columns.get_level_values(0):
            prices[ticker] = raw_data[ticker]['Close']

# Clean the data
print("Cleaning data...")
# Remove columns with all NaN values
prices = prices.dropna(axis=1, how='all')
# Remove rows with all NaN values
prices = prices.dropna(how='all')
# Forward fill missing values
prices = prices.fillna(method='ffill')
# Backward fill remaining NaN values
prices = prices.fillna(method='bfill')

final_tickers = prices.columns.tolist()
print(f" Clean price data: {prices.shape[0]} dates × " +
      f"{prices.shape[1]} tickers")
print(f"Date range: {prices.index[0]} to {prices.index[-1]}")
print(f"Final tickers count: {len(final_tickers)}")

# Display uniformly formatted output
print("Displaying first 5 rows of the cleaned price data:")
print(prices.head(5))

# Plot the first 5 tickers to visualize the data
print("Plotting first 5 tickers...")
prices.iloc[:, :5].plot(figsize=(14, 7))
plt.title('Adjusted Close Prices of Top 5 S&P 500 Tickers')
plt.xlabel('Date')
plt.ylabel('Adjusted Close Price')

```

```
plt.legend(loc='upper left')
plt.grid()
plt.show()
```

Extracting and cleaning adjusted close prices...

Cleaning data...

Clean price data: 1254 dates × 100 tickers

Date range: 2020-08-07 00:00:00 to 2025-08-05 00:00:00

Final tickers count: 100

Displaying first 5 rows of the cleaned price data:


	AAPL	GOOG	GOOGL	AMZN	AVGO	\
Date						
2020-08-07	108.203766	74.282944	74.471870	158.373001	28.915100	
2020-08-10	109.776505	74.362968	74.394836	157.408005	29.041964	
2020-08-11	106.511765	73.578644	73.585678	154.033493	28.746538	
2020-08-12	110.051590	74.885864	74.912727	158.112000	29.599096	
2020-08-13	111.999245	75.473877	75.380417	158.050995	29.224720	

	BRK-B	COST	ABBV	BAC	CVX	...	\
Date						...	
2020-08-07	209.479996	314.168182	75.978493	23.083565	69.710876	...	
2020-08-10	212.580002	313.329498	75.536949	23.481403	72.064003	...	
2020-08-11	212.660004	306.353363	75.774078	23.799673	71.975677	...	
2020-08-12	213.240005	310.343781	78.096283	23.631697	72.859085	...	
2020-08-13	211.979996	309.366821	77.417595	23.295742	72.136269	...	

	ACGL	A	BR	BRO	DXCM	\
Date						
2020-08-07	30.894695	94.532845	124.429207	44.604279	110.175003	
2020-08-10	31.122908	93.914055	125.671173	44.478378	105.305000	
2020-08-11	31.237015	93.101952	130.740326	44.216911	102.787498	
2020-08-12	30.999290	94.387802	127.989563	44.352482	106.972504	
2020-08-13	30.771076	95.489929	127.731949	44.371853	108.820000	

	STZ	AWK	AEE	ADM	AVB
Date					
2020-08-07	158.571671	136.478409	71.071022	38.364456	129.755493
2020-08-10	158.506912	135.858871	71.996376	38.765251	129.704926
2020-08-11	159.819748	130.243774	70.188896	38.678131	127.967621
2020-08-12	162.251221	133.398849	70.898056	38.730404	128.954315
2020-08-13	163.286636	133.984116	70.370522	38.712975	127.090576

[5 rows x 100 columns]  
Plotting first 5 tickers...



#### Task 4: Compute Daily Log Returns

Formula:  $r_t = \ln(P_t / P_{t-1}) = \ln(P_t) - \ln(P_{t-1})$

```
# Compute daily log returns:  $r_t = \ln(P_t / P_{t-1})$ 
print("Computing daily log returns...")

log_returns = np.log(prices / prices.shift(1))
log_returns = log_returns.dropna()

# Clean any infinite or NaN values
log_returns.replace([np.inf, -np.inf], np.nan, inplace=True)
log_returns.fillna(method='ffill', inplace=True)
log_returns.fillna(0, inplace=True)

print(f" Log returns computed successfully!")
print(f"Shape: {log_returns.shape}")
print(f>Date range: {log_returns.index[0]} to {log_returns.index[-1]}")

# Basic statistics
print(f"\nSummary statistics:")
```

```

print(f"Mean daily return: {log_returns.mean().mean():.6f}")
print(f"Std daily return: {log_returns.std().mean():.6f}")
print(f>Data quality - NaN values: {log_returns.isnull().sum().sum()}")
print(f>Data quality - Infinite values: {np.isinf(log_returns).sum().sum()}")

# Display first 5 rows of log returns
print("Displaying first 5 rows of log returns:")
print(log_returns.head(5))

# Plot log returns: cleaner and more informative visualization
import matplotlib.dates as mdates

plt.figure(figsize=(16, 7))
# Plot only a subset of tickers for clarity (e.g., top 10 by market cap)
top10 = top_100_tickers[:10]
log_returns[top10].plot(ax=plt.gca(), alpha=0.7, linewidth=1.2)

plt.title('Daily Log Returns of Top 10 S&P 500 Stocks by Market Cap',
          fontsize=18, fontweight='bold')
plt.xlabel('Date', fontsize=14)
plt.ylabel('Log Return', fontsize=14)
plt.legend(top10, loc='upper right', ncol=2, fontsize=10,
           frameon=True)
plt.grid(True, linestyle='--', alpha=0.5)
plt.tight_layout()
plt.xlim(log_returns.index[0], log_returns.index[-1])
plt.gca().xaxis.set_major_locator(mdates.YearLocator())
plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%Y'))
plt.show()

```

Computing daily log returns...

Log returns computed successfully!

Shape: (1253, 100)

Date range: 2020-08-10 00:00:00 to 2025-08-05 00:00:00

Summary statistics:

Mean daily return: 0.000565

Std daily return: 0.019336

Data quality - NaN values: 0

Data quality - Infinite values: 0

Displaying first 5 rows of log returns:

AAPL	GOOG	GOOGL	AMZN	AVGO	BRK-B	\
------	------	-------	------	------	-------	---



Date


2020-08-10	0.014430	0.001077	-0.001035	-0.006112	0.004378	0.014690
2020-08-11	-0.030191	-0.010603	-0.010936	-0.021671	-0.010224	0.000376
2020-08-12	0.032694	0.017610	0.017873	0.026134	0.029226	0.002724
2020-08-13	0.017543	0.007821	0.006224	-0.000386	-0.012729	-0.005926
2020-08-14	-0.000892	-0.007085	-0.007957	-0.004121	-0.004869	-0.004823

	COST	ABBV	BAC	CVX	...	ACGL	A \
Date					...		
2020-08-10	-0.002673	-0.005828	0.017088	0.033198	...	0.007360	-0.006567
2020-08-11	-0.022516	0.003134	0.013463	-0.001226	...	0.003660	-0.008685
2020-08-12	0.012941	0.030186	-0.007083	0.012199	...	-0.007639	0.013717
2020-08-13	-0.003153	-0.008728	-0.014318	-0.009970	...	-0.007389	0.011609
2020-08-14	0.001726	0.004111	0.004544	0.005884	...	-0.003405	-0.014276

	BR	BRD	DXCM	STZ	AWK	AEE \
Date						
2020-08-10	0.009932	-0.002827	-0.045209	-0.000408	-0.004550	0.012936
2020-08-11	0.039544	-0.005896	-0.024197	0.008248	-0.042209	-0.025426
2020-08-12	-0.021264	0.003061	0.039908	0.015099	0.023936	0.010053
2020-08-13	-0.002015	0.000437	0.017123	0.006361	0.004378	-0.007469
2020-08-14	0.000288	0.000000	-0.023524	-0.003346	-0.008224	-0.001230

	ADM	AVB
Date		
2020-08-10	0.010393	-0.000390
2020-08-11	-0.002250	-0.013485
2020-08-12	0.001351	0.007681
2020-08-13	-0.000450	-0.014558
2020-08-14	0.003595	0.015473

[5 rows x 100 columns]



## Deliverables

```
# DELIVERABLE 1: prices - Adjusted close price matrix (DataFrame)
print("DELIVERABLE 1: prices")
print("=" * 40)
print(f"Type: {type(prices)}")
print(f"Shape: {prices.shape}")
print(f"Index: {prices.index[0]} to {prices.index[-1]}")
print(f"Columns: {len(prices.columns)} tickers")
print("\nSample data (first 5 rows, first 5 columns):")
display(prices.iloc[:5, :5])

print("\n" + "=" * 40)

# DELIVERABLE 2: log_returns - Daily log returns matrix (DataFrame)
print("DELIVERABLE 2: log_returns")
print("=" * 40)
print(f"Type: {type(log_returns)}")
print(f"Shape: {log_returns.shape}")
print(f"Index: {log_returns.index[0]} to {log_returns.index[-1]}")
print(f"Columns: {len(log_returns.columns)} tickers")
print(f"Formula:  $r_t = \ln(P_t / P_{t-1})$ ")
print("\nSample data (first 5 rows, first 5 columns):")
display(log_returns.iloc[:5, :5])
```

```

print("\n" + "=" * 40)

# DELIVERABLE 3: This notebook - Data preprocessing script
print("DELIVERABLE 3: This notebook")
print("=" * 40)
print("Type: Jupyter Notebook (.ipynb)")
print("Contains all code and comments for data acquisition " +
      "and preprocessing")

```

DELIVERABLE 1: prices

```

=====
Type: <class 'pandas.core.frame.DataFrame'>
Shape: (1254, 100)
Index: 2020-08-07 00:00:00 to 2025-08-05 00:00:00
Columns: 100 tickers

```

Sample data (first 5 rows, first 5 columns):

	AAPL	GOOG	GOOGL	AMZN	AVGO
Date					
2020-08-07	108.203766	74.282944	74.471870	158.373001	28.915100
2020-08-10	109.776505	74.362968	74.394836	157.408005	29.041964
2020-08-11	106.511765	73.578644	73.585678	154.033493	28.746538
2020-08-12	110.051590	74.885864	74.912727	158.112000	29.599096
2020-08-13	111.999245	75.473877	75.380417	158.050995	29.224720

```

=====
DELIVERABLE 2: log_returns
=====
Type: <class 'pandas.core.frame.DataFrame'>
Shape: (1253, 100)
Index: 2020-08-10 00:00:00 to 2025-08-05 00:00:00
Columns: 100 tickers
Formula: r_t = ln(P_t / P_{t-1})

```

Sample data (first 5 rows, first 5 columns):

	AAPL	GOOG	GOOGL	AMZN	AVGO
Date					
2020-08-10	0.014430	0.001077	-0.001035	-0.006112	0.004378
2020-08-11	-0.030191	-0.010603	-0.010936	-0.021671	-0.010224
2020-08-12	0.032694	0.017610	0.017873	0.026134	0.029226
2020-08-13	0.017543	0.007821	0.006224	-0.000386	-0.012729
2020-08-14	-0.000892	-0.007085	-0.007957	-0.004121	-0.004869

=====

DELIVERABLE 3: This notebook

=====

Type: Jupyter Notebook (.ipynb)

Contains all code and comments for data acquisition and preprocessing